

# **Neural Nets & Theano Python Library**

# **1: Neural Nets**

# Neural Nets

- Linear combination of features

$$b + w_1 * x_1 + w_2 * x_2 + \dots + w_p * x_p$$

- Non-linear activation (e.g.  $\mathbf{wx} \rightarrow [-1, 1]$ )

$$f(b + w_1 * x_1 + w_2 * x_2 + \dots + w_p * x_p)$$



- Pass outputs on as inputs (multi-layer nets)
  - This gives non-linear decision boundaries
  - This gives a non-convex problem :(

# Neural Nets


- Linear combination of features

$$b + w_1 * x_1 + w_2 * x_2 + \dots + w_p * x_p$$


- Non-linear activation




- Non-linear activation:
  - Once linear combination exceeds a cutoff dramatic change in behavior




- Non-linear activation:
  - Once linear combination exceeds a cutoff dramatic change in behavior



Q: why use  
nonlinear  
activation instead  
of a *step activation*  
(no transition)?



Q: why use  
nonlinear  
activation instead  
of *linear*  
*activation*?






# Example: Logistic Regression

Linear combination of features (log-odds),  
passed into sigmoid activation function to get  $p$ .


$$\begin{aligned} \text{log odds} &= \log \frac{p}{1-p} = x_1\beta_1 + \dots + x_m\beta_m \\ \implies p &= \exp(\text{log odds}) / (1 + \exp(\text{log odds})) \end{aligned}$$

# Layering for nonlinearity

- Single layer nets (e.g. logistic regression) can make linear decisions:



Linearly separable (you can draw a straight line to separate these two colors)




Not linearly separable (you can't draw a straight line to separate the two colors)

# Layering for nonlinearity

- Multi-layer nets can approximate any function, i.e. can do any nonlinear separation.
  - Downside: the optimization problem is nonconvex
  - Downside: you may need *lots* of nodes / layers

# Neural Net OR

**bias**  
(intercept) **= -15**




**w**

X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	1
?	?	

# Neural Net OR

**bias**  
(intercept) **= -15**




**w**

X1	X2	WX + bias	Y
0	0	-15	~0
0	1	5	~1
1	0	5	~1
1	1	25	~1
+20	+20		

Either X  
being  
active  
enough  
to turn  
on Y

# Neural Net AND

**bias**  
(intercept) **= -15**




**w**

X1	X2	Y
0	0	0
0	1	0
1	0	0
1	1	1
?	?	

# Neural Net AND

**bias**  
(intercept) **= -15**



**w**


X1	X2	WX + bias	Y
0	0	-15	~0
0	1	-5	~0
1	0	-5	~0
1	1	5	~1
+10	+10		

Need both X to be active to turn on Y

# Neural Net XOR

*can't do it with one layer*

**bias**  
(intercept) **= -15**




**w**

X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0
?	?	



# Neural Net XOR

*insert layer that does AND,  
OR (we know how to make  
AND and OR operations!)*




X1	X2	X1& X2	X1  X2	Y
0	0	0	0	0
0	1	0	1	1
1	0	0	1	1
1	1	1	1	0
		?	?	

# Neural Net XOR

*assume we have outputs of  
AND and OR layer...*

**bias**  
(intercept) **= -15**




**w**

A1:= X1&X2	A2:= X1 X2	Y
0	0	0
0	1	1
0	1	1
1	1	0
?	?	

# Neural Net XOR

*assume we have outputs of  
AND and OR layer...*

**bias**  
(intercept) **= -15**



**w**

A1:= X1&X 2	A2:= X1 X2	wA + bias	Y
0	0	-15	~0
0	1	5	~1
0	1	5	~1
1	1	-5	~0
-10	+20		

The OR  
alone  
can turn  
Y on, but  
the AND  
disables  
it.

# Neural net computation

- In practice: many layers, many nodes, very non-convex
  - Lots of calculations
    - Many nodes
    - Try lots of starting values
  - Parallel computing (most nodes don't interact with each other)
    - e.g. for XOR, could compute AND, OR in parallel

# Backpropagation briefly

- Perceptron:
  - Find error, update parameters so error might get fixed
- Back-propagation:
  - Neural nets fit using gradient descent
  - BP also finds an error, at the output layer, then it sends it backwards through the neural net

## **2: Theano**

# Theano

- It looks weird, why? Basically, speed:
  - Python (numpy) is not the most efficient language
  - Python is not parallelized
- Fitting neural nets:
  - Lots of computation
  - Better to be parallel (e.g. use a GPU / map-reduce)

# Aside: why GPUs?

- CPU: central processing unit.
  - Does program execution
- GPU: graphics processor unit.
  - Does calculations to render things for display on screen



# Aside: why GPUs?

- CPU: central processing unit.
  - Few cores, each fast & “smart”, good @ serial tasks
- GPU: graphics processor unit.
  - Many cores, each slow & “dumb”, good @ parallel tasks

# Aside: why GPUs?

- CPU: central processing unit.
  - Few cores, each fast & “smart”, serial
    - Smart: more features, e.g. you can run an OS
- GPU: graphics processor unit.
  - Many cores, each slow & “dumb”, parallel
    - Dumb: it can only do certain tasks well (linear algebra)

# Aside: why GPUs?

- CPU: central processing unit.
  - Few cores, each fast & “smart”, serial
- GPU: graphics processor unit.
  - Many cores, each slow & “dumb”, parallel
  - **Better for neural nets: lots of simple, parallel calculations (in sum, faster than the CPU)**

# Theano

- Allows construction of “more efficient code”
  - Alternatives: inline a bunch of C code
- Theano can talk to a GPU
  - Recall: GPUs are “dumb”, you usually have to write another language to use them
- For both cases, theano acts like a “foreign language interpreter”

# Theano Code: *Symbolic Representation*

```
>>> x = T.dmatrix('x')
```

```
>>> s = 1 / (1 + T.exp(-x))
```

```
>>> logistic = function([x], s)
```

```
>>> logistic([[0, 1], [-1, -2]])
```

```
array([[ 0.5, 0.73105858],  
       [ 0.26894142, 0.11920292]])
```

# Theano Code: *Symbolic Representation*

```
>>> x = T.dmatrix('x')
```

```
>>> s = 1 / (1 + T.exp(-x))
```

```
>>> logistic = function([x], s)
```

I'm confused! Why do we define  $x$  and  $s$  this way? Where's the data?

# Theano Code: *Symbolic Representation*

```
>>> x = T.dmatrix('x')    Define input type
```

```
>>> s = 1 / (1 + T.exp(-x)) Define function behavior
```

```
>>> logistic = function([x], s)
```

Put it all together

# Theano Code: *Optimization (1)*

```
w = theano.shared(np.asarray((np.random.randn(*(numFeatures,  
numClasses))*0.01)))
```

Weights

```
X = T.matrix()
```

```
Y = T.matrix()
```

Data



# Theano Code: *Optimization (1)*

```
w = theano.shared(np.asarray((np.random.randn(*(numFeatures,  
numClasses))*.01)))
```

Weights

Shared: can update

```
X = T.matrix()
```

Data

```
Y = T.matrix()
```

Tensor: cannot update

# Theano Code: *Optimization (2)*

```
def model(X, w):  
    return T.nnet.softmax(T.dot(X, w))
```

Loss (“cost”)

```
y_hat = model(X, w)
```

```
cost = T.mean(T.nnet.categorical_crossentropy(y_hat, Y))
```

```
gradient = T.grad(cost=cost, wrt=w)
```

```
update = [[w, w - gradient * alpha]]
```

Gradient  
Descent

# Theano Code: *Optimization (2)*

```
def model(X, w):  
    return T.nnet.softmax(T.dot(X, w))
```

```
y_hat = model(X, w)
```

```
cost = T.mean(T.nnet.categorical_crossentropy(y_hat, Y))
```

```
gradient = T.grad(cost=cost, wrt=w)
```

```
update = [[w, w - gradient * alpha]]
```

Loss (“cost”)

Theano “knows” many  
functions

Theano “knows” the gradient  
of these functions

Gradient  
Descent

## Theano Code: *Optimization* (3)

```
train = theano.function(inputs=[X, Y],  
                        outputs=cost,  
                        updates=update)
```

Theano 'function' set up:

1. Set inputs (data)
2. Set output (current loss; "cost")
3. Set update rule (do gradient descent)

# Things to try besides the notebook:

Use theano to write:

An AND neural net

An OR neural net

An XOR net

# Things to try besides the notebook:

AND, OR rough guide:

1. Copy the logistic regression code
2. What is the data?
  - a. Two binary features
  - b. Binary output
3. Don't run gradient descent too long!
4. Instead of checking accuracy, see what weights you get.